

EX.NO:4
DATE:28/01/2025

VALUE ITERATION ALGORITHM

AIM:

To implement the Value Iteration algorithm to solve the Frozen Lake problem in OpenAI Gym, calculate the optimal policy, and determine the optimal state-value function. The goal is to navigate the environment effectively by maximizing cumulative rewards while avoiding hazards.

PROCEDURE:

1. Environment Initialization:

Import the necessary libraries (numpy and gym).

Create the Frozen Lake environment using gym.make(). Ensure is_slippery=True for a stochastic environment.

2. Define Value Iteration:

Initialize variables:

- n_states: Total number of states in the environment.
- n_actions: Total number of possible actions.
- V: State-value function initialized to zeros.
- directions: Symbols ('<', 'v', '>', '^') to represent actions visually.

Set up parameters:

- gamma: Discount factor for rewards.
- theta: Threshold for convergence.

3. Perform Iterative Updates:

Use a while loop to repeatedly update the state-value function $V(s)$:

- For each state s:
 - Calculate the action-value function $Q(s,a)$ for all possible actions a using the Bellman equation: $Q(s,a) = \sum_{s'} P(s'|s,a) \cdot [R(s,a,s') + \gamma \cdot V(s')]$
 - Update $V(s)$ with the maximum $Q(s,a)$.
 - Record the action with the highest $Q(s,a)$ for visualizing the policy directions.
- Track the maximum change in $V(s)$ to check for convergence.

4. Extract Optimal Policy:

After convergence, extract the optimal policy:

- For each state, find the action a that maximizes $Q(s,a)$.
- Store the action in the policy array.
- Update a policy matrix with symbols to visually represent the optimal actions.

CODE:

```
import numpy as np
import gym

def value_iteration(env, gamma=0.99, theta=1e-8):

    n_states = env.observation_space.n
```

```

n_actions = env.action_space.n

V = np.zeros(n_states) # Initialize state-value function

policy = np.zeros(n_states, dtype=int) # Initialize policy

while True:

    delta = 0

    for s in range(n_states):

        q_values = np.zeros(n_actions)

        for a in range(n_actions):

            for prob, next_state, reward, done in env.P[s][a]:

                q_values[a] += prob * (reward + gamma * V[next_state])

        max_value = np.max(q_values)

        delta = max(delta, abs(max_value - V[s]))

        V[s] = max_value

    if delta < theta:

        break

    for s in range(n_states):

        q_values = np.zeros(n_actions)

        for a in range(n_actions):

            for prob, next_state, reward, done in env.P[s][a]:

                q_values[a] += prob * (reward + gamma * V[next_state])

        policy[s] = np.argmax(q_values)

return policy, V

# Run Value Iteration on Frozen Lake

env = gym.make("FrozenLake-v1", is_slippery=True)

optimal_policy, optimal_values = value_iteration(env)

```

```
print("Optimal Policy:")  
  
print(optimal_policy)  
  
print("Optimal State-Value Function:")  
  
print(optimal_values)
```

OUTPUT:

```
Final Optimal Policy Directions:  
[['<' '^' '^' '^']  
 ['<' '<' '<' '<']  
 ['^' 'v' '<' '<']  
 ['<' '>' 'v' '<']]  
  
Optimal Policy:  
[0 3 3 3 0 0 0 0 3 1 0 0 0 2 1 0]  
  
Optimal State-Value Function:  
[0.54202581 0.49880303 0.47069551 0.4568515  0.55845085 0.  
 0.35834799 0.          0.59179866 0.64307976 0.6152075  0.  
 0.          0.7417204  0.86283741 0.          ]
```